

Image processing using OpenCV + Python + mjpg-streamer

Dustin Spicuzza (Team 2423/1418)

September 10, 2016

NE FIRST University Day

Agenda

- Why OpenCV + Python?
- Image filtering demo
- pynetworktables

Image processing

- FRC Teams do it a lot of ways
 - NIVision (LabVIEW)
 - GRIP (Uses OpenCV as engine)
 - OpenCV (various custom stuff)
- We're going to talk about OpenCV

Why OpenCV?

- Originally developed by Intel
- It has thousands of image processing related algorithms and functions available
 - Highly optimized and reliable
 - Has building blocks that fit together
- Lets you do complex image processing without needing to understand the math
 - If you understand the math, it helps!

Why OpenCV?

- Bindings for multiple languages
 - C/C++
 - Java
 - Python
- Multiple platforms supported
 - Windows
 - Linux
 - OSX
 - Android
- Oh, and it's **FREE!**

What OpenCV Provides

- Image I/O:
 - Read/Write images from disk
 - Use native OS functionality to interface with cameras
- Image Segmentation
 - Edge finding
 - Contour detection
 - Thresholding

What OpenCV Provides

- Face detection
- Motion tracking
- Stereo vision support
- Support for GPU acceleration
- Machine learning operations
 - Classifiers
 - Neural networks

What OpenCV Provides

- Distributed with lots of useful samples that you can use to figure out how OpenCV works
 - Face detection
 - Edge finding
 - Histograms
 - Square finder

Lots and lots and lots of stuff...

Why Python + OpenCV?

- Python is really easy to learn and use
 - Simple syntax
 - Rapid prototyping
- Most of the compute intensive work is implemented in C/C++
 - Python is just glue, realtime operation **is** possible
- NumPy is awesome
 - Manipulating image data is trivial compared to other OpenCV bindings (Java, C++)

Time to CODE!

Go to <http://goo.gl/nB0NCG>

About this environment

<http://goo.gl/nB0NCG>

- It's a Jupyter Notebook (formerly IPython Notebook)
 - This slideshow uses Jupyter too!
- It allows you to mix text and executable code in a webpage
- You execute each cell using SHIFT-ENTER

Hello World!

- Click the cell with the following text, and press SHIFT-ENTER

In [2]: `print("Hello class")`

Hello class

Next Steps

- Execute the helper code
- The next cell tells you about the images available in your environment

In [3]:


```
%ls images
```

2013-f0.png	2016-cmp-0.jpg	2016-cmp-5.jpg	2016-p0.jpg
2013-p0.png	2016-cmp-1.jpg	2016-dcmp1.jpg	2016-p1.jpg
2013-p1.png	2016-cmp-3.jpg	2016-dcmp2.jpg	2016-p2.jpg
2014-f0.png	2016-cmp-4.jpg	2016-dcmp3.jpg	2016-p3.jpg

Hello image!

- Let's load an image and show it


```
In [4]: # Change this to load different images  
img = cv2.imread('images/2016-cmp-5.jpg')  
imshow(img)
```



Hello image!

- You can show multiple images next to each other

In [5]: `imshow(img, img)`



OpenCV Image Basics

- Images are stored as multidimensional arrays
 - Color images have 3 dimensions: height, width, channel
- Each pixel is a number stored in the array
- Numpy array notation allows you to do operations on individual pixels or ranges of pixels

```
In [6]: img[50, 150, :]          # Access a single pixel,
```

```
Out[6]: array([ 8, 19,  0], dtype=uint8)
```

```
In [7]: x = img[24:42, 42:100, :]    # Access a range of pixels
```



OpenCV Image Basics

- Color is represented by storing combinations of Red, Blue, and Green pixels in separate channels
 - OpenCV uses BGR representation, not RGB
- The amount of each individual color is represented in the individual channel
 - 'dark' is zero, 'bright' is 255
- Combine the channels to represent a color
 - Green = RGB(0, 255, 0)
 - Deep Pink = RGB(255, 20, 147)

OpenCV Image Basics

- Using numpy we can easily fill an image with a single color

```
In [8]: # define image with height=240, width=320, 3 channels  
shape = (240, 320, 3)  
pink_img = np.empty(shape, dtype=np.uint8)  
  
# Fill every pixel with a single color  
pink_img[:] = (0, 255, 0)  
  
imshow(pink_img)
```



Practical Example

- 2016 FIRST Stronghold: find targets that are surrounded by retroreflective tape, and shoot boulders into them

Practical Example

- Finding gray tape at a distance isn't particularly easy
 - Key part of image processing is removing as much non-essential information from image
- We can do better!

Retroreflective Tape

- It has a useful property -- it reflects light directly back at the source
- What can we do with this property?
- Shine bright LEDs at the target and the tape reflects that color back to the camera
 - Many teams have found that green light works best
- Reduce exposure of camera so only bright light sources are seen

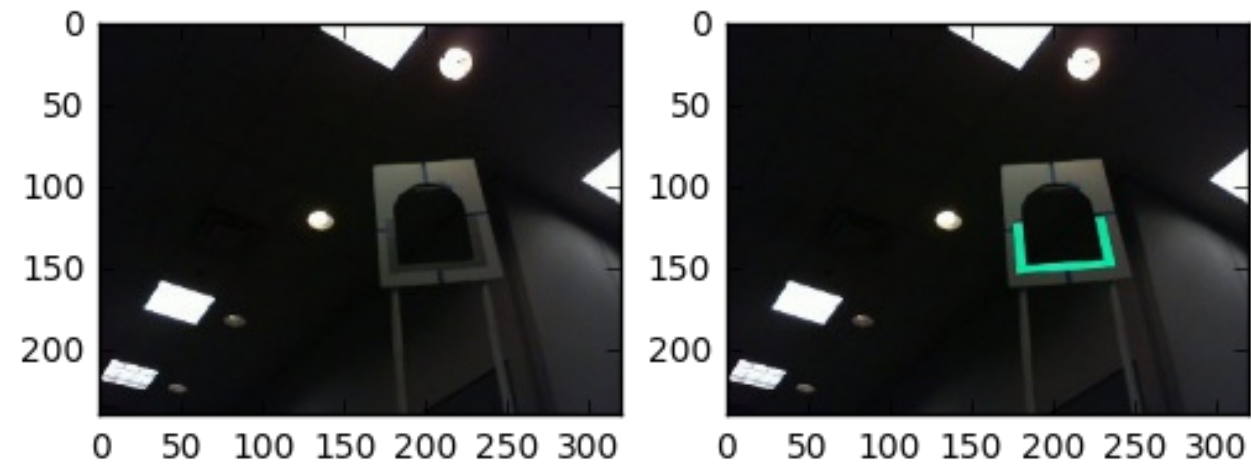
A note about exposure

- Webcams support setting the exposure manually (yay)
- Some cameras only allow particular exposure settings
 - The lifecam is one of them
- OpenCV has bugs, it doesn't set the exposure properly
- Here's a workaround that works on linux:

```
v4l2-ctl -d /dev/video0 -c exposure_auto=1 -c exposure_absolute=10
```

Retroreflective Tape

```
In [9]: img1 = cv2.imread('images/2016-p0.jpg')  
img2 = cv2.imread('images/2016-p1.jpg')  
imshow(img1, img2)
```



Practical Example

Processing steps to find targets:

- Isolate the green portions of the image
- Analyze the green portions to determine targets

Note: There are a lot of ways to go about this, I'm just showing you one way


Identify the green

- What is “green” anyways?
 - This is green. This is also green.
- To a computer, green is really a range of colors
- An object’s color changes depending on lighting conditions
- We can transform the image to identify colors independent of lighting conditions

Identify the green

- Convert the image from RGB to HSV
 - Hue: the color
 - Saturation: Colorfulness
 - Value: Brightness


```
In [10]: hsv = cv2.cvtColor(img, cv2.COLOR_BGR2HSV)  
imshow(img, hsv)
```



Identify the green

That doesn't show why HSV is useful. Let's look at the individual channels instead.

```
In [11]: h, s, v = cv2.split(hsv)
         imshow(h, s, v)
```



Identify the green


- Green is a range of values present in the image
- 'Threshold' the image to get rid of the colors that we don't care about
- Lots of ways to do this
 - Manually specify values
 - Automated methods

Identify the green

`cv2.inRange` can threshold an image given two ranges of pixels.

- Wanted values are converted to 255
- Unwanted values are now 0

```
In [12]: lower = np.array([0, 145, 80])  
         upper = np.array([255, 255, 255])  
  
         filtered = cv2.inRange(hsv, lower, upper)  
         imshow(img, filtered)
```




Identify the green

Sometimes, you end up with holes in your output

```
In [13]: img3 = cv2.imread('images/2013-f0.png')
hsv3 = cv2.cvtColor(img3, cv2.COLOR_BGR2HSV)


# Thresholds are different because different camera/lighting
lower3 = np.array([30, 188, 16])
upper3 = np.array([75, 255, 255])
filtered3 = cv2.inRange(hsv3, lower3, upper3)
imshow(filtered3)
```



Identify the green

- We can use a morphological operation to fill in the holes
 - Various types of morphology operations available
- They modify a pixel based on the values of its neighboring pixels
 - The one we use to fill in holes is a “closing” operation

```
In [14]: kernel = cv2.getStructuringElement(cv2.MORPH_RECT, (2,2), anchor=(1,1))
output = cv2.morphologyEx(filtered3, cv2.MORPH_CLOSE, kernel,
                                     iterations=9)
imshow(output)
```



Identifying Targets

Use `findContours()` to find regions of interest


- Returns a list of points bounding each separate blob in the image (called a contour)
- Also returns a hierarchy so you can determine whether a contour is entirely inside another contour

[illegible]

Identifying Targets

If you want to see what it found, you can draw the found contours.

```
In [16]: dst = np.zeros(shape=img.shape, dtype=img.dtype)
cv2.drawContours(dst, contours, -1, (0, 255, 255), 1)
imshow(dst)
#print(contours[0])
```



Identifying Targets

- As you can see, contours aren't the whole story

Identifying Targets

- Contour analysis
 - Discard non-convex contours
 - Convert to polygon approximation (approxPolyDP)
 - Discard polygons that aren't rectangles
 - Discard polygons that aren't the right size

Magic?

```
In [17]: min_width = 20 # in pixels
results = []

# Iterate over each contour
for c in contours:

    # Contours are jagged lines -- smooth it out using an approximation
    a1 = cv2.approxPolyDP(c, 0.01*cv2.arcLength(c, True), True)


    # This fills in the contour so that it's a rectangle
    hull = cv2.convexHull(c)

    # Approximate the points again, smoothing out the hull
    a2 = cv2.approxPolyDP(hull, 0.01*cv2.arcLength(hull, True), True)

    # We only care about objects that are wider than they are tall, and things w
ider
    # than a particular width. Only keep things that meet that criteria.
    _, _, w, h = cv2.boundingRect(a2)
    if w > h and w > min_width and len(a2) in (4,5):
        results.append(a2)
```

Magic?

```
In [18]: # Finally, draw out our results  
dst = np.zeros(shape=img.shape, dtype=img.dtype)  
cv2.drawContours(dst, results, -1, (0, 0, 255), 1)  
imshow(dst)  
#print(results[0])
```



Identifying Targets

- Sometimes you need to do more work
 - Use ratios to determine which target you're looking at
 - Remove duplicates (inner rectangles)
 - Other types of validation

Now what?

We have targets... probably should do something with them?

Calculate angle/distance to target

- I'm not a math guy, but this sorta works
 - Angle works, distance is a bit iffy
- Get the minimum bounding rectangle
- Figure out the horizontal and vertical field of view for your camera
 - Look it up online
- Do math to it

In [19]: *# Just do the first one for now*
result = results[0]

Get the height/width
h = float(img.shape[0])
w = float(img.shape[1])

Define HFOV and VFOV
VFOV = 45.6 *# degrees*
HFOV = 61.0 *# degrees*

In [20]: `((cx, cy), (rw, rh), rotation) = cv2.minAreaRect(result)`

`# These work fairly well`

`angle = VF0V * cy / h - (VF0V/2.0)`

`height = HF0V * cx / w - (HF0V/2.0)`

`print(angle, height)`

`-15.564704704284669 3.2886390209198026`

```
In [21]: # This is magic, but it doesn't really work  
target_height = 7.66 # 7' 8"  
camera_height = 1.08 # 13"  
camera_pitch = 40.0 # What angle is the camera at?  
t = (target_height - camera_height)  
distance = t/math.tan(math.radians(-angle + camera_pitch))  
  
print(distance)
```

4.5113763224432235

Now What?

- Send data via NetworkTables
- ... I forgot to write this slide. It's easy, I promise.

Where to run the image processing

- RoboRIO
 - RoboRIO is relatively slow, OpenCV eats a lot of CPU
 - Hint: Make the images small (320x240)
 - Less hardware to deal with
 - FIRST intends to install OpenCV by default in 2017

Where to run the image processing

- Driver Station
 - Streaming images to OpenCV is possible
 - Various latency bugs
 - Latency is an issue here
 - mDNS problems (hopefully will be resolved in 2017)

Where to run the image processing

- Coprocessor (Jetson, Raspberry PI, Nexus 5)
 - Lots of teams do this
 - More hardware to deal with
 - Potentially higher fidelity processing

mjpg-streamer

- Open Source application for streaming Webcam over HTTP
- Allows accessing camera stream from a webpage
- Very little CPU usage if no image processing happening

mjpg-streamer

- I created an OpenCV plugin for it
- You can provide python processing code:

```
class MyFilter:
    def process(self, img):
        pass

def init_filter():
    f = MyFilter()
    return f.process
```

mjpg-streamer

- I would show you a demo here, but my RoboRIO's USB is toasted

mjpg-streamer

- Peter Johnson is working on a better FRC-specific version of this
- Expect to see it in 2017

Want code?

- Working OpenCV code integrated with mjpg-streamer
 - <https://github.com/frc2423/2016/tree/master/OpenCV>
 - Includes code for storing images onto USB drive during matches
 - Don't let our robot's performance fool you... :(
- The stuff we did here will be available sometime tonight
 - <https://github.com/virtuald/frc-imageprocessing-workshop-2016>

If you want more

- Team 254 gave an excellent presentation at CMP in 2016
 - <https://goo.gl/mppi4E>
 - Video/audio:
<http://www.chiefdelphi.com/forums/showthread.php?t=147568&page=3>
 - Latency compensation is an excellent technique presented here

Resources

- Python 3.5.x
 - <https://www.python.org/downloads/>
- Learn Python
 - <http://www.codecademy.com/tracks/python>
- OpenCV 3.1.0
 - <http://opencv.org>
- NumPy
 - Official site: <http://www.numpy.org>

Resources

- roborio-packages
 - <https://github.com/robotpy/roborio-packages>
- OpenCV for RoboRIO
 - <https://github.com/robotpy/roborio-opencv>
- mjpg-streamer for RoboRIO
 - <https://github.com/robotpy/mjpg-streamer>

Resources

- pynetworktables
 - source code + examples @ <https://github.com/robotpy/pynetworktables>
- Edit & debug python code using Eclipse
 - Pydev: <http://pydev.org/>

One more thing...

FIRSTwiki: <https://firstwiki.github.io>

- Publicly editable repository of information related to FIRST Robotics
 - Technical topics
 - Non-technical
 - Team pages
- Add content to your team's page!

Questions?